

Approximate and Exact Geometric Generalized Minimum Spanning Trees

Majid Mirzanezhad 

Electrical Engineering and Computer Science, Ohio University, USA

Arash Rafiey 

Indiana State University, IN, USA

Abstract

We revisit the geometric variant of the generalized minimum spanning tree (GGMST) problem: we are given a set of n points in \mathbb{R}^d inside an integer grid with N non-empty clusters where each cluster is a unit hypercube, the aim is to find a minimum spanning tree that contains exactly one point from each cluster while respecting the adjacencies of the clusters. For any $\epsilon > 0$, we give an $(1 + \epsilon)$ approximation algorithm whose running time is $O^*(N(\frac{1}{\epsilon})^{O(\frac{1}{\epsilon})} + N^{d+1} + n \log n)$, where O^* hides factors polynomial in $1/\epsilon$. We also give an exact FPT algorithm that runs in $O(N^2 \rho^2 (2^{O(\kappa^4)}))$, where κ is an implicit parameter that relates to the input integer grid containing the non-empty clusters. Here, ρ is the maximum number of points across all clusters. When a GGMST instance is given in the plane (\mathbb{R}^2), we present $(1 + \epsilon)$ -approximation with running time $O(N\rho^2 280^{(\frac{1}{\epsilon})})$. In the case that for any non-empty cluster, either the row or column it belongs to, consists of at most $\gamma < N$ many non-empty clusters, we give a faster FPT algorithm that runs in $O(N\rho^2 560^\gamma)$ time.

2012 ACM Subject Classification CCS

Keywords and phrases Computational Geometry, Minimum Spanning Tree, Approximation Algorithm

Digital Object Identifier 10.4230/LIPIcs.CVIT.2016.23

Funding Arash Rafiey : Supported by Bailey College Faculty Fellowship

1 Introduction

In *Generalized Minimum Spanning Tree* (GMST) problem, we are given an undirected abstract graph $G = (V, E)$ with vertex set V and edge set $E \subseteq V \times V$ together with a weight function $w : E \rightarrow \mathbb{R}^+$ and a partition of vertices V into non-intersecting clusters. The goal is to find a minimum weight tree containing exactly one vertex from each cluster. This problem is a generalization of the well-known *Minimum Spanning Tree* (MST) Problem where each vertex of the graph is a cluster itself. In the *Generalized Traveling Salesman Problem* (GTSP) we are given a graph $G = (V, E)$ such that V is a set of vertices and $E \subseteq V \times V$. Moreover, a weight function $w : E \rightarrow \mathbb{R}^+$ assigns weights to the edges and the set of vertices V is partitioned into non-intersecting clusters. The goal is to find a minimum weight cycle that goes through all clusters and from each cluster, there exists exactly one point in that cycle.

In the geometric variant of the problem above, we are given N non-empty unit d -dimensional axis-aligned unit hypercubes, which we refer to as *clusters*, and a set of n points in \mathbb{R}^d spread across the N clusters. Each cluster contains at least two points. The aim is to pick exactly one point from each cluster and connect the chosen N points by straight-line segments (Euclidean edges) to form a minimum spanning tree (MST). We only allow edges between points chosen from *adjacent* clusters, where adjacency means the clusters share at least a corner. This problem is also known as *Geometric Generalized Minimum Spanning Tree* (GGMST). We refer to this problem as GGMST problem.



© Majid Mirzanezhad and Arash Rafiey;
licensed under Creative Commons License CC-BY 4.0
42nd Conference on Very Important Topics (CVIT 2016).

Editors: M.Mirzanezhad and A. Rafiey; Article No. 23; pp. 23:1–23:16

 Leibniz International Proceedings in Informatics

Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

44 **1.1 Related work**

45 The GGMST problem occurs in telecommunications network planning, where a network of
 46 node clusters need to be connected via a tree architecture using exactly one node per cluster [7].
 47 More precisely, local subnetworks must be interconnected by a global network containing a
 48 gateway from each subnetwork. For this inter-networking, a point has to be chosen in each
 49 local network as a hub and the hub point must be connected via transmission links such as
 50 optical fiber [13]. Additionally, the problem has some applications in design of backbones in
 51 large communication networks, energy distribution, and agricultural irrigation [9].

52 The GMST problem was first introduced by Myung et al. (1995) [13]. However, the
 53 conventional MST problem is polynomially solvable [6], it was shown in [13] that the GMST
 54 problem is strongly NP-hard and there is no constant factor approximation for it unless
 55 $P = NP$. It is worth mentioning that some heuristic algorithms have been developed for
 56 GMST. For more information, see [7, 9].

57 When approximation is allowed, Pop et al. [14] lay out a 2ρ -approximation algorithm
 58 for GGMST where the size of each cluster is bounded by ρ . They take advantage of Integer
 59 programming and LP relaxation to tackle the problem. Feremans, Grigoriev, and Sitters
 60 consider GGMST [5] and show a strong NP-hardness, even if we restrict to instances in which
 61 all non-empty clusters are connected and each cluster contains at most two points (Theorem
 62 1 in [5]). They design a dynamic programming algorithm which solves the problem with
 63 connected non-empty clusters in time $O(l\rho^{6k}2^{34k^2}k^2)$ where the grid has size $l \times k$ and it is
 64 polynomial when k is fixed (Theorem 2 in [5]). Moreover, authors in [5] suggest a polynomial
 65 time approximation scheme (PTAS) for GGMST problem when the non-empty grids are
 66 connected. It is easy to see that their dynamic programming algorithm is highly exponential
 67 in k , as a result their PTAS using the dynamic program may not be perceived as efficient
 68 in practice. To overcome this issue, Bhattacharya et al. designed a simple polynomial-time
 69 approximation algorithm for GGMST with a guaranteed ratio of $1 + 4\sqrt{2} + \epsilon$ [4].

70 Later, Kachooei et al. [10] showed that the problem is still NP-hard when there are at
 71 most two points in each cluster even with equal y -coordinates, and it is unlikely to have
 72 a fully polynomial time approximation scheme (FPTAS) for it unless $P = NP$. It is also
 73 known that the GGMST is NP-hard even if the cluster graph forms a tree [15].

74 **1.2 Our contribution**

75 We obtained two primary results in higher dimensions and then we show how they can be
 76 improved when the instance is given in the plane. In particular, we propose approximation
 77 and FPT algorithms for each case as follows:

- 78 1. We give a $(1 + \epsilon)$ -approximation algorithm that runs in $O(n \log n + n/\epsilon^d + (N/\epsilon)^{d+1} +$
 79 $N(\frac{1}{\epsilon})^{O(1/\epsilon)})$, in any constant dimension d (Thm. 8). Our approach utilizes the idea of
 80 Arora's original PTAS technique for TSP in [2]. Arora's technique requires portal *pairings*
 81 (entering and exiting a quadtree cell in matching pairs) to form a tour. Our GGMST
 82 variant tracks *connected* subsets of portals (partitions) to encode how GGMST edges
 83 unite boundary portals to form connected components, i.e., forests within each cell of
 84 the quadtree constructed over the clusters. In TSP, every point must appear in the tour,
 85 however in our case the problem specifies exactly one point must be chosen per cluster.
 86 We encode these portal "activations" at quadtree leaf cells via a dynamic program and
 87 avoid re-choosing points at higher levels of the quadtree.
- 88 2. We also give an exact FPT algorithm that runs in $O(\rho^2 N^2(2^{O(k^4)}))$, where κ is the
 89 treewidth of the hypergraph induced by the non-empty clusters of the integer grid in \mathbb{R}^d ,

- 90 for any constant dimension d (Thm. 12). In telecom networks, especially large-scale ones,
 91 each local subnetwork (e.g., LAN, building, or regional cluster) designates exactly one node
 92 (its gateway or hub) to connect to external links. This is precisely the setting where each
 93 cluster (local subnetwork) has to select one hub. Real networks often have “hierarchical”
 94 or “near-tree” structures for reasons of reliability, scalability, or cost. For instance, metro
 95 fiber networks or networks built on geographic constraints (where expensive cross-links
 96 are limited) may indeed exhibit sparse topological connections, effectively bounding the
 97 treewidth.
- 98 3. We show that the optimal GGMST in \mathbb{R}^2 can be computed in time $O(N\rho^2 280^k)$ where
 99 the instance $(\mathcal{G}, \mathcal{N}, P)$ lies in a grid of size $k \times l$ and ρ is the maximum number of
 100 points inside each cluster (Thm. 14). As a byproduct of this result, we can obtain the
 101 following PTAS. Using a similar approximation argument that appeared in [5], one can
 102 obtain a $(1 + \epsilon)$ -approximation for instance $(\mathcal{G}, \mathcal{N}, P)$ in time $O(N\rho^2 280^{\frac{1}{\epsilon}})$ (Corollary 15)
 103 which significantly improves the $(1 + \epsilon)$ -approximation algorithm with running time
 104 $O\left(\frac{l}{\epsilon^2}\rho^{\frac{6}{\epsilon}}2^{\frac{34}{\epsilon^2}}\right)$ in [5]. Note that our approximation algorithm in the plane can outperform
 105 compared to the one that works in higher dimensions if the quality of approximation is
 106 of interest since the base is a constant number compared to $1/\epsilon$.
- 107 4. In the case that for any non-empty cluster, either the row or column it belongs to, consists
 108 of at most $\gamma < N$ many non-empty clusters, we give another FPT algorithm in \mathbb{R}^2 that
 109 runs in $O(N\rho^2 560^\gamma)$ which is much faster than our FPT algorithm in \mathbb{R}^d (Thm. 19).

110 2 Preliminaries

111 Given two points $x, y \in \mathbb{R}^d$, the Euclidean distance between them is denoted by $\|x - y\|$,
 112 which represents the standard ℓ_2 norm, capturing the natural geometric separation in the
 113 space. Computing minimum spanning trees requires setting a fixed point as the root of a tree.
 114 We denote a subtree rooted at a point x by T_x . This notation emphasizes that T_x consists of
 115 all nodes and edges descending from x , encapsulating the hierarchical structure induced by
 116 the root. The weight of the tree $W(T)$ represents the total edge lengths of T . When dealing
 117 with substructures, we use $W(T_x)$ to specify the weight of the subtree rooted at x . These
 118 definitions serve as the foundation for analyzing spatial structures when optimizing the
 119 weights overall possible GGMSTs.

120 ▶ **Definition 1** (Hypergraph (HG)). A “Hypergraph” of an integer grid is a graph whose
 121 vertices are the non-empty clusters in the grid and two vertices share an edge iff the clusters
 122 are adjacent.

123 Let $\mathcal{I} = (\mathcal{G}, \mathcal{N}, P)$ be an instance of the GGMST where \mathcal{G} is an HG, \mathcal{N} is the set of
 124 clusters with $|\mathcal{N}| = N$, and P is the set of points spread across the clusters in \mathcal{N} , with
 125 $|P| = n$. An HG \mathcal{G} in \mathcal{I} can attain a spanning tree whose nodes span all the clusters in \mathcal{I} .
 126 The said spanning tree is defined as follows:

127 ▶ **Definition 2** (Backbone Spanning Tree (BBST)). A “Backbone Spanning Tree (BBST)”
 128 w.r.t. some integer grid clusters is a spanning tree if
 129 1. Each vertex corresponds to a cluster,
 130 2. All vertices span the entire non-empty clusters,
 131 3. And two vertices share an edge if their corresponding clusters are adjacent to each other.

132 ▶ **Definition 3.** Let $\mathcal{I} = (\mathcal{G}, \mathcal{N}, P)$ be an instance of the GGMST problem. For a given BBST
 133 T , let $\text{OptMST}(T)$ denote the optimal tree corresponding to the BBST T .

23:4 Approximate and Exact Geometric Generalized Minimum Spanning Trees

134 Note that every backbone spanning tree can induce a GGMST whose adjacency of nodes
135 follows the adjacency of nodes in the BBST. Let \mathcal{B} be the set of all possible BBSTs. Then:

$$\text{136} \quad T_{opt} := \min_{T \in \mathcal{B}} \text{OptMST}(T).$$

137 In other words, a branching tree of the BBST T describes the adjacency of the vertices of a
138 GGMST T^* where T^* selects the points of every cluster marked by T and connects them
139 with respect to the adjacency of the clusters in T .

140 3 Approximation algorithm for GGMST

141 As mentioned earlier, the GGMST problem is NP-hard in general, but some underlying
142 geometric structural properties can be utilized to achieve approximation schemes. Our goal
143 is to provide a PTAS, for every $\epsilon > 0$, albeit potentially with a high dependence on $1/\epsilon$.

144 3.1 Algorithmic sketch

145 We begin with discretizing each cluster into a finite set of candidate points independent of n .
146 By choosing a sufficiently fine grid inside each cluster, we can guarantee that picking the
147 nearest candidate point to the optimal solution's chosen point will not increase the weight
148 (W) of optimal GGMST by more than $\epsilon \cdot W(T_{opt})$. Next, we construct a *WSPD-based* $(1 + \epsilon)$ -
149 *spanner* to reduce the potential $O(n^2)$ edge set down to $O(n)$ edges without significantly
150 increasing GGMST cost.

151 We then build a *quadtree* on the clusters, placing $O(1/\epsilon)$ *portals* on each cell boundary to
152 constrain how GGMST edges may cross those boundaries. Applying a bottom-up *dynamic*
153 *programming* (DP) merges partial GGMST solutions from child cells by matching their
154 portal-connectivity patterns similar to Arora's technique for TSP that appeared in [2]. By
155 carefully bounding the "overhead", i.e., inflated weight added to the solution due to the
156 approximation via the portal placement introduced at each scale (from forcing edges to go
157 through portals) and leveraging the spanner's limited set of edges, the algorithm achieves a
158 total GGMST cost within $(1 + \epsilon)$ of the true optimum, all in polynomial time for a fixed ϵ .

159 Let $\delta > 0$, we partition each cluster uniformly into grid cells each of side length δ . This
160 yields $O(\frac{1}{\delta^d})$ candidate points per cluster. Let P be the resulting set of all cluster grid points.
161 The naïve algorithm below is a dynamic programming algorithm that can compute the exact
162 GGMST given a BBST T .

163 ▶ **Lemma 4.** *For a given BBST T , one can compute the $\text{OptMST}(T)$ in polynomial time.*

164 **Proof.** Select an internal vertex of T corresponding to a cluster C as the root. For any leaf
165 cluster $l \in T$, and for any point $p \in l$, set $W(T_p) = 0$. For any internal node $I \in T$, and any
166 point $q \in I$, let $\text{Chd}(I)$ denote the set of children of I in BBST T . Then we have:

$$\text{167} \quad W(T_q) = \sum_{c \in \text{Chd}(I)} \min_{p \in c} \left\{ W(T_p) + \|p - q\| \right\}$$

168 The algorithm runs in $O(N\rho^2)$ time using a forward and then backward breadth-first search
169 on T since the number of children per node is $O(2^{d+1}) = O(1)$, where $d = O(1)$ and $\rho < n$ is
170 the maximum number of points in a cluster. ◀

171 ▶ **Lemma 5 (Discretization).** *Let $\epsilon > 0$. We can choose $\delta = \epsilon / (\sqrt{d}(N - 1))$ sufficiently small
172 so that there exists a $(1 + \epsilon)$ -approximate GGMST solution selecting from these candidate
173 points. Moreover, the number of grid points overall is $O(N^d / \epsilon^d)$*

¹⁷⁴ **Proof.** Consider an optimal GGMST solution T_{opt} that selects a point q_i in each cluster C_i .
¹⁷⁵ By construction of the grid inside C_i , there is a grid point p_i no more than $\delta\sqrt{d}$ away from
¹⁷⁶ q_i (since the cluster is unit size and we have a δ -spaced lattice).

¹⁷⁷ Replace each q_i by p_i to obtain a new (approximate) solution T' . The increase in the
¹⁷⁸ length of any edge (q_i, q_j) in T_{opt} is at most proportional to δ because shifting points can
¹⁷⁹ increase distances by at most $\sqrt{d}\delta$. Thus, the total additional cost incurred by this rounding
¹⁸⁰ is at most $\sqrt{d}\delta N$, since the MST has at most $N - 1$ edges and each edge length changes by
¹⁸¹ at most $\sqrt{d}\delta$.

¹⁸² In the case that $N \leq 2^d$ and assuming $d = O(1)$, one can enumerate all $N^{N-2} = 2^{(d2^d)} =$
¹⁸³ $O(1)$ many BBSTs and compute GGMST with a minimum total weight in polynomial time
¹⁸⁴ in an exact manner utilizing the algorithm described in Lemma 4. For all $N > 2^d$, following
¹⁸⁵ the packing argument demonstrated in Lemma 1 in [3], we derive that $W(T_{\text{opt}}) > 1$. Hence:

$$\begin{aligned} \suparrow 186 \quad W(T') &\leq W(T_{\text{opt}}) + (N - 1) \cdot \sqrt{d}\delta = W(T_{\text{opt}}) + (N - 1) \cdot \frac{\epsilon}{N-1} = W(T_{\text{opt}}) + \epsilon \leq \\ \suparrow 187 \quad W(T_{\text{opt}}) + \epsilon \cdot W(T_{\text{opt}}) &= (1 + \epsilon) \cdot W(T_{\text{opt}}). \end{aligned} \quad \blacktriangleleft$$

¹⁸⁸ 3.2 The dynamic program

¹⁸⁹ We present the algorithm and analysis in four main steps: (1) quadtree construction and
¹⁹⁰ cluster isolation, (2) building a $(1 + \epsilon)$ -Spanner using WSPD, (3) portal placement and the
¹⁹¹ dynamic programming scheme, and finally, (4) analysis of approximation guarantee and
¹⁹² running Time.

¹⁹³ The first two steps are relatively straightforward to handle as follows: our first goal
¹⁹⁴ is to spatially organize the clusters and points so that we can apply a dynamic program
¹⁹⁵ to propagate the minimum weight of a spanning forest from each “zone” to the adjacent
¹⁹⁶ ones. We use a *quadtree* to recursively partition the plane until each leaf cell of the quadtree
¹⁹⁷ contains at most one cluster (thus fully containing that cluster). By ensuring each cluster
¹⁹⁸ resides in precisely one leaf cell, picking a representative grid point from that cluster is a
¹⁹⁹ local decision at the leaf. We utilize further to ensure a constant number of points per leaf
²⁰⁰ limits state explosion in the DP state.

²⁰¹ Without further techniques, considering an MST among n points might require looking at
²⁰² $O(n^2)$ potential edges. To make the problem tractable, we reduce complexity by constructing
²⁰³ a $(1 + \epsilon)$ -spanner. Given n points in the plane and $\epsilon > 0$, it is known that one can
²⁰⁴ construct a $(1 + \epsilon)$ -spanner G with $O(n/\epsilon^d)$ edges in $O(n \log n + n/\epsilon^d)$ time [8]. Furthermore:
²⁰⁵ $\text{MST}(G) \leq (1 + \epsilon)\text{MST}(P)$, where P is the set of points and $\text{MST}(P)$ is the optimal MST
²⁰⁶ cost on these n points.

²⁰⁷ 3.2.1 Portal placement and the DP scheme

²⁰⁸ Even with a spanner, we face the challenge of ensuring we pick exactly one point per cluster
²⁰⁹ to form a GGMST. We approach this as follows: In the Arora-style PTAS algorithm, portals
²¹⁰ are used to discretize how solutions cross cell boundaries. Place $p = O(1/\epsilon)$ portals uniformly
²¹¹ along each side of every quadtree cell. Any MST edge crossing the boundary must “snap”
²¹² to the nearest portal. By introducing portals: i) we limit the complexity of boundary
²¹³ interactions in the DP states and ii) we ensure that any additional cost from snapping edges
²¹⁴ to portals is at most $\epsilon \cdot W(T_{\text{opt}})$ after appropriate scaling takes effect.

²¹⁵ ▶ **Lemma 6.** Let T_{opt} be an GGMST of cost $W(T_{\text{opt}})$ in a bounding box of side-length M .
²¹⁶ Suppose at quadtree level i , each cell has side length $\frac{M}{2^i}$. Let X_i be the number of GGMST
²¹⁷ edges that cross cell boundaries at level i . If every crossing edge is at least $\alpha \frac{M}{2^i}$ in length,

23:6 Approximate and Exact Geometric Generalized Minimum Spanning Trees

218 then:

$$219 \quad X_i \leq \frac{\beta W(T_{opt})}{\alpha (M/2^i)}$$

220 for some constant $\alpha > 0$ and $\beta > 0$.

221 ► **Lemma 7** (Portal Approximation). *There is a way to choose $O(1/\epsilon)$ portals per cell boundary so that rerouting GGMST edges through these portals increases the GGMST cost by at most $\epsilon \cdot W(T_{opt})$.*

224 **Proof.** Following the geometric scaling argument from Arora's PTAS, at each level of the
225 quadtree, the cost added by snapping edges to portals is a small fraction of the GGMST's
226 cost *at that scale*. Summing these fractions over all levels (cells at different scales) yields at
227 most $\epsilon \cdot W(T_{opt})$ in total overhead.

228 Let the bounding box of the entire cluster set have side length M . The quadtree is built by
229 recursively splitting into four equal squares. At level 0 (the root), cells have side length M .
230 At level i , cells have side length $\frac{M}{2^i}$. We place $p = O(\frac{1}{\epsilon})$ portals on each cell boundary, so
231 the portal spacing at level i is on the order of $\frac{M}{2^i \cdot p} = O(\epsilon \frac{M}{2^i})$. The cost analysis unfolds as
232 follows:

233 *Overhead per Crossing at Level i :* Consider an GGMST edge e that crosses a boundary
234 at level i . Snapping e from its original crossing point to the nearest portal adds at most
235 $O(\epsilon \frac{M}{2^i})$ extra length, since the portal spacing is $O(\epsilon \frac{M}{2^i})$.

236 *Bounding the Number of Crossings:* We now focus on all GGMST edges crossing cell
237 boundaries at level i . Each such crossing indicates an edge of length at least $\Theta(\frac{M}{2^i})$, since it
238 spans across different cells at that scale. Following Lemma 6, if there were too many such
239 "long" edges, their combined length would exceed $W(T_{opt})$, contradicting the GGMST's
240 minimality. Hence, we conclude that the *total length* of edges crossing boundaries at level i
241 is bounded by some constant factor of $W(T_{opt})$. Since each crossing can add up to $O(\epsilon \frac{M}{2^i})$
242 extra length, the total *overhead* at this level is at most $(\epsilon \cdot c) \cdot W(T_{opt})$ for some universal
243 constant $c > 0$.

244 *Summation Over All Levels:* The quadtree may have $O(\log N)$ levels (or $O(\log M)$). If we
245 allocated a full $\epsilon W(T_{opt})$ overhead *per level*, we could end up with $O(\log N) \cdot \epsilon W(T_{opt})$. To
246 avoid this, we use a *geometric distribution* of overhead:

$$247 \quad \text{Overhead}_i \leq \frac{\epsilon}{2^i} W(T_{opt}) \quad \text{at each level } i.$$

248 By placing portals more densely at coarser levels, the detour at scale i is forced to be only
249 $\frac{\epsilon}{2^i} W(T_{opt})$. Summing these over $i = 0$ to $\log N$ gives

$$250 \quad \text{Total Overhead} = \sum_{i=0}^{\log N} \text{Overhead}_i \leq \epsilon W(T_{opt}) \sum_{i=0}^{\log N} \frac{1}{2^i} < 2\epsilon W(T_{opt}).$$

251 Adjusting constants or using $\epsilon' = \frac{\epsilon}{2}$ in the construction ensures the total overhead does not
252 exceed $\epsilon \cdot W(T_{opt})$. This completes the proof. ◀

253 3.2.2 DP state definition and encoding connectivity

254 Consider a cell C in the quadtree. The DP state $D(C, S, B)$ might encode:

255 1. C : The cell for which a GGMST is being computed.

- 256 2. S : Which clusters fully contained in C have chosen their representative point. This
 257 controls the GGMST to be a forest and then connects to other forests at a higher-level
 258 cell. Since each leaf cell has at most one cluster, and these sets merge as we go up, S is
 259 manageable.
- 260 3. B : A description of how GGMST connections that leave C through its boundary portals
 261 are arranged. This is a pattern representing how portals are connected internally.

262 **Portal Partitions:** To represent how portals connect inside C , we consider a partition of
 263 the set of portals on C 's boundary into connected components. Each connected component
 264 represents a partial GGMST structure inside C that connects some subset of these boundary
 265 portals. The number of partitions of a set of m elements is given by the m th Bell number
 266 B_m , which grows super-exponentially in k . For $m = p = O(1/\epsilon)$, we know that Bell number
 267 B_m is:

$$268 \quad B_m = \Theta((m^m / (e^m \cdot m^{3/2}))).$$

269 Hence $B_m \leq m^m = O(1/\epsilon)^{O(1/\epsilon)}$. This encoding is crucial. By storing a “minimal structure”
 270 that indicates which portals are connected, we can later merge these configurations at
 271 higher-level cells.

272 **DP Recurrence:** For a leaf cell C , we know exactly which points it contains (at most a
 273 constant number). We try all possible ways of picking one point for the cluster (if there is
 274 one) and then determine how this affects portal connectivity. In particular, for each choice,
 275 we compute the partial GGMST cost inside C (often trivial since it is just one cluster plus
 276 some boundary edges to portals). Then, we encode the resulting portal connectivity in a
 277 partition B .

$$278 \quad D(C, S, B) = \underbrace{\text{PortalsCost}(p, B)}_{\text{how } p \text{ connects to boundary portals}}$$

279 where $\text{PortalsCost}(p, B)$ is the cost of hooking up p to the boundary portals in a way that
 280 yields the partition B . In other words, how does p connect to (some or all) of the portals in
 281 B ? This ensures that the internal MST structure is consistent with B . For an internal cell
 282 C with children C_1, C_2, C_3, C_4 , we combine their solutions:

$$283 \quad D(C, S, B) = \min_{\substack{S_1 \cup S_2 \cup S_3 \cup S_4 = S \\ B_1, B_2, B_3, B_4}} \left(\sum_{i=1}^4 D(C_i, S_i, B_i) + \text{MERGE COST}(B_1, B_2, B_3, B_4, B) \right),$$

284 where MERGE COST involves adding edges from the spanner G between portals to ensure
 285 the combined configuration at C respects partition B . The spanner's $O(n)$ edges ensure
 286 checking merges is efficient, see Figure 1.

287 3.2.3 Analysis: approximation and running time

288 The approximation factors set to $\epsilon/3$ for various steps such as cluster discretization, spanner
 289 construction, and portal partitioning combine to yield a $(1 + \epsilon)$ -approximation overall. We
 290 now give a detailed view of the running time analysis, including how the complexity of
 291 handling portals is controlled. The runtime analysis is as follows:

- 292 1. We discretize each of N cell into $O(N^d/\epsilon^d)$ cells, hence it takes $O(N^{d+1}/\epsilon^d)$ time.
 293 2. We build a $(1 + \epsilon/3)$ -spanner G in time of $O(n \log n + n/\epsilon^d)$.

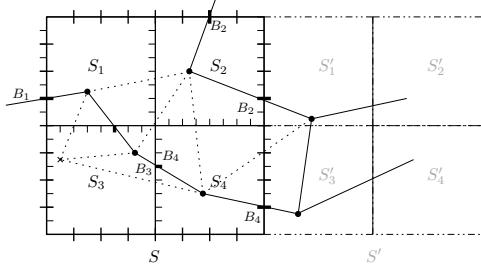


Figure 1 Connected components in the cell $S = S_1 \cup S_2 \cup S_3 \cup S_4$ have formed through a detour through the adjacent cell S' reckoned by MERGE COST. The dotted lines display the edges of the underlying spanner. A partition of active (bold) portals on the boundary can be stored in B .

- 294 3. We build a quadtree on N cells in time $O(N \log N)$. The quadtree has $O(N)$ cells since
 295 we stop subdividing when each cell has at most a constant number of points and at most
 296 one cluster. Each leaf cell is thus simple to handle. Suppose each leaf has $l = O(\frac{N^d}{\epsilon^d})$
 297 candidate points. For each candidate grid point, we check edges to up to $p = O(\frac{1}{\epsilon})$
 298 portals. That's $O(l \cdot k) = O(\frac{N^d}{\epsilon^d} \cdot \frac{1}{\epsilon})$. Summation across all leaves, since we have $O(N)$
 299 leaves, the total is $N \times O(\frac{N^d}{\epsilon^d} \cdot \frac{1}{\epsilon}) = O(\frac{N^{d+1}}{\epsilon^{d+1}})$.
- 300 4. At internal cells, we use a dynamic programming (DP) approach. The DP states consider:
 301 – Which of the four subcells are represented in the cell (chosen at lower levels).
 302 – How the GGMST inside the subcells connects to one another outside via a set of
 303 boundary portals. The challenge is that at each cell, we must consider different ways to
 304 connect portals. The number of ways to partition a set of m elements into connected
 305 subsets is at most the m th Bell number $B_m = (1/\epsilon)^{O(1/\epsilon)}$. Since the number of the
 306 cells in the quadtree is $O(N)$, the total would be $O(N(1/\epsilon)^{O(1/\epsilon)})$.
 307 Overall, the runtime is $O(N^{d+1}/\epsilon^d) + O(n \log n + n/\epsilon^d) + O(N \log N) + O((N/\epsilon)^{d+1}) + O(N \cdot$
 308 $(1/\epsilon)^{O(1/\epsilon)})$, which is: $O(n \log n + n/\epsilon^d + (N/\epsilon)^{d+1} + N \cdot (1/\epsilon)^{O(1/\epsilon)})$.

309 ▶ **Theorem 8.** *Given an instance $\mathcal{I} = (\mathcal{G}, \mathcal{N}, P)$ with $|\mathcal{N}| = N$ and $|P| = n$, where clusters
 310 are axis-aligned unit hypersquares in \mathbb{R}^d , for any $\epsilon > 0$, one can compute a $(1+\epsilon)-\text{GGMST}$
 311 of \mathcal{I} in time $O^*((1/\epsilon)^{O(1/\epsilon)} N + N^{d+1} + n \log n)$, where O^* hides factors polynomial in $1/\epsilon$.*

312 4 An FPT algorithm for exact GGMST

313 Our algorithm breaks down into two levels by mixing a coarse dynamic program to enumerate
 314 all BBSTs like T and a fine-grained dynamic programming algorithm to obtain $\text{OptMST}(T)$
 315 (using the algorithm in Lemma 4). However, the number of backbone spanning trees can grow
 316 exponentially with the size of the graph, making enumeration computationally challenging.
 317 Given an instance $\mathcal{I} = (\mathcal{G}, \mathcal{N}, P)$ with $|\mathcal{N}| = N$ and $|P| = n$, we present an algorithm
 318 that efficiently enumerates all spanning trees of our $\mathcal{G} = (V, E)$ with bounded treewidth κ ,
 319 exploiting the structure provided by a nice tree decomposition, where V is the set of clusters
 320 and E is the set of edges representing the adjacency between the hypernode in V . Computing
 321 the exact treewidth of a graph is NP-complete [1]. However, Korhonen [12] presented an
 322 algorithm that computes a tree decomposition of width k' satisfying $\kappa' \leq 2\kappa + 1$, where κ is
 323 the optimal treewidth of \mathcal{G} .

324 ▶ **Theorem 9** (Korhonen's Algorithm [12]). *Given a graph G with treewidth κ , there exists an
 325 algorithm that computes a tree decomposition of \mathcal{G} with width $k' \leq 2\kappa + 1$ in time $O(2^{O(k)} \cdot N)$.*

326 In our algorithm, we use the tree decomposition obtained from Korhonen's algorithm
 327 and κ' will be used as the treewidth parameter in our analysis.

328 ▶ **Definition 10** (Nice Tree Decomposition). *A tree decomposition is nice if:*

329 1. *It is a rooted tree.*

330 2. *Each node is of one of the following types:* Leaf: *A leaf node with an empty bag.* Introduce:
 331 *A node with one child t' such that $X_t = X_{t'} \cup \{v\}$ for some $v \notin X_{t'}$.* Forget: *A node with*
 332 *one child t' such that $X_t = X_{t'} \setminus \{v\}$ for some $v \in X_{t'}$.* Join: *A node with two children t_1*
 333 *and t_2 such that $X_t = X_{t_1} = X_{t_2}$.*

334 Every tree decomposition can be transformed into a nice tree decomposition with at most
 335 $O(N)$ nodes [11].

336 Our goal is to enumerate all combinatorially different spanning trees of \mathcal{G} . We achieve
 337 this by performing dynamic programming over the nice tree decomposition of \mathcal{G} . The key
 338 idea is to represent partial solutions at each node of the decomposition and combine them
 339 according to specific rules that ensure correctness and avoid redundancy.

340 4.1 DP formulation

341 Our algorithm consists of several components that update the state of connectivity of vertices
 342 in each bag and the acyclicity of the spanning tree induced by the vertices in the bag.

343 *State Representation:* At each node t of the nice tree decomposition, we maintain a set of
 344 DP states. Each state represents a set of partial spanning trees consistent with the subgraph
 345 induced by the vertices seen so far.

346 ▶ **Definition 11** (DP State). *A DP state at node t is a tuple (σ, δ) where:*

347 1. *σ is a partition of the vertices in X_t , representing the connected components among them*
 348 *in the partial spanning tree.*

349 2. *$\delta \subseteq E_t$ is the set of edges included in the partial spanning tree, where E_t is the set of*
 350 *edges between vertices in X_t .*

351 Note that each state must satisfy the following invariants: (1) *Acyclicity:* The edges in δ
 352 do not form cycles, and (2) *Consistency:* The partition σ accurately reflects the connectivity
 353 induced by δ . Let DP_t denote the set of DP states at node t . We define recursive formulas
 354 for computing DP_t based on the type of node. We have the following breakdowns for our
 355 DP recurrences:

356 ■ *Leaf Nodes:* $DP_t = \{(\emptyset, \emptyset)\}$.

357 ■ *Introduce Nodes:* Let t be an introduce node introducing vertex v , with child t' . The DP
 358 states are computed as:

$$359 DP_t = \bigcup_{(\sigma', \delta') \in DP_{t'}} \bigcup_{\delta_v \subseteq E_v} \{(\sigma_v, \delta_v \cup \delta') \mid \text{Valid}(\sigma_v, \delta_v \cup \delta')\}$$

360 where: E_v is the set of edges between v and vertices in $X_{t'}$, δ_v is a subset of E_v (possible
 361 ways v can connect to $X_{t'}$), σ_v is the updated partition obtained by adding v to σ' and
 362 merging parts if v is connected to vertices in δ_v , and $\text{Valid}(\sigma_v, \delta_v \cup \delta')$ ensures acyclicity
 363 and consistency.

364 ■ *Forget Nodes:* Let t be a forget node forgetting vertex v , with child t' . The DP states are
 365 computed as:

$$366 DP_t = \{(\sigma|_{X_t}, \delta') \mid (\sigma', \delta') \in DP_{t'}\}$$

367 where $\sigma|_{X_t}$ is the restriction of the partition σ' to X_t .

23:10 Approximate and Exact Geometric Generalized Minimum Spanning Trees

368 ■ *Join Nodes:* Let t be a join node with children t_1 and t_2 . The DP states are computed as:

$$369 \quad \text{DP}_t = \{(\sigma, \delta_1 \cup \delta_2) \mid (\sigma, \delta_1) \in \text{DP}_{t_1}, (\sigma, \delta_2) \in \text{DP}_{t_2}, \text{Valid}(\sigma, \delta_1 \cup \delta_2)\}$$

370 The function $\text{Valid}(\sigma, \delta)$ returns true if: The edges in δ do not form cycles. The partition
371 σ correctly represents the connected components induced by δ . See Appendix 7 for further
372 explanation and detailed pseudocode describing the recursive formulas above.

373 4.2 Runtime and space analysis

374 Let $N = |V|$ be the number of vertices in \mathcal{G} and κ' be the width of the tree decomposition
375 obtained from Korhonen's algorithm. At each node t , the size of the bag is $|X_t| \leq \kappa' + 1$.
376 The number of possible partitions σ of X_t is given by the Bell number $B_{\kappa'+1}$ and the number
377 of possible edge subsets δ among vertices in X_t is $2^{\binom{\kappa'+1}{2}}$. Since for any $m > 0$, the m th Bell
378 number is: $B_m = \left((m/e)^m m^{3/2} \right)$, the total number of states at each node, and hence the
379 required space is:

$$380 \quad S(N, \kappa) = B_{\kappa'+1} \cdot 2^{\binom{\kappa'+1}{2}} = O\left(2^{O(\kappa^2)} \left(O(\kappa/e)^{2\kappa+2} O(\kappa)^{3/2}\right)\right) = O\left(2^{O(k^2)}\right),$$

381 Since $\kappa' \leq 2\kappa + 1$. Regarding time complexity, we need to explore each:

- 382 ■ *Introduce Nodes:* For each state in $\text{DP}_{t'}$, we consider $2^{|E_v|}$ subsets δ_v . Note that $|E_v| \leq \kappa'$,
383 thus we have $2^{\kappa'}$ possibilities. The runtime per introduce node: $O(S \cdot 2^{\kappa'}) = O(S \cdot 2^{O(\kappa)})$.
- 384 ■ *Forget Nodes:* For each state in $\text{DP}_{t'}$, we perform operations in $O(1)$ time. Hence The
385 runtime per forget node: $O(S)$.
- 386 ■ *Join Nodes:* For each pair of states from the two children, we check for compatibility and
387 merge. Number of state pairs is $O(S^2)$.

388 To obtain the total runtime: The number of nodes in the nice tree decomposition is $O(N)$.
389 Hence, the total time complexity is:

$$390 \quad T(N, \kappa) = O\left(N(S^2 + S \cdot 2^{O(\kappa)})\right) = O\left(N\left(2^{O(\kappa^4)}\right)\right)$$

391 ▶ **Theorem 12.** Given an instance $\mathcal{I} = (\mathcal{G}, \mathcal{N}, P)$ with $|\mathcal{N}| = N$ and $|P| = n$, where clusters
392 are axis-aligned unit hypersquares in \mathbb{R}^d , one can compute the exact GGMST of \mathcal{I} in time
393 $O\left(\rho^2 N^2 \left(2^{O(\kappa^4)}\right)\right)$, where ρ is the maximum number of points in across all clusters in \mathcal{N} .

394 **Proof.** As shown above, given an instance \mathcal{I} , with κ as the treewidth of \mathcal{G} , one can enumerate
395 all possible BBSTs in time $T(N, \kappa) = O\left(N\left(2^{O(\kappa^4)}\right)\right)$. Once all BBSTs like T are obtained,
396 we can explicitly pick points from the selected clusters that result in the optimal GGMST,
397 i.e., $\text{OptMST}(T)$ using Lemma 4. During enumeration of all BBSTs, we fix each $T \in \text{BBST}$
398 and utilize the algorithm described in Lemma 4 that runs in $O(N\rho^2)$. Therefore, the total
399 runtime would be: $T(N, n, \kappa) = O\left(\rho^2 N^2 \left(2^{O(\kappa^4)}\right)\right)$. ◀

400 5 Approximation and FPT algorithms for GGMST in the plane

401 Let $(\mathcal{G}, \mathcal{N}, P)$ be an example of the GGMST problem with N non-empty cells where the
402 non-empty cells are connected. Let ρ be the maximum number of points inside each cluster.

403 ▶ **Lemma 13.** The number of backbone forests in a grid $k \times 3$ (with some non-empty cells
404 that) is at most 280^k .

405 **Proof.** Let $F(k)$ denote the number of such forests. Let C_1, C_2, C_3 denote the clusters in
 406 row $k - 1$ and C_4, C_5, C_6 be the clusters in row k . Let H be a graph with vertices C_1, \dots, C_6
 407 and edges $(C_1, C_4), (C_1, C_5), (C_2, C_4), (C_2, C_5), (C_2, C_6), (C_3, C_5), (C_3, C_6), (C_4, C_5), (C_5, C_6)$.
 408 Here, the set of edges of H is the adjacency of each of the clusters C_1, \dots, C_6 without
 409 considering the adjacency between C_1, C_2 and C_2, C_3 cells. $F(k) \leq h \times F(k - 1)$ where h
 410 is the number of forests that are sub-graphs of G . It is not difficult to design a computer
 411 simulation and verify that $h = 280$. Therefore, $F(k) \leq 280^k$. ◀

412 ▶ **Theorem 14.** *Let $\mathcal{I} = (\mathcal{G}, \mathcal{N}, P)$ be an instance of GGMST with $|\mathcal{N}| = N$ and $|P| = n$
 413 and with k fixed rows. The optimal GGMST of \mathcal{I} can be computed in $O(N\rho^2 280^k)$ time.*

414 **Proof.** Let $\mathcal{F}(i, j)$, where $i \leq j$, denote the set of possible forests restricted to the columns
 415 from i to j , covering all k rows of \mathcal{G} . By Lemma 13, there exist at most 280^k trees from $\mathcal{F}(1, 3)$
 416 that can serve as the backbone of \mathcal{G} . For each such tree T , we apply the dynamic programming
 417 approach described in Lemma 4 to obtain $\text{OptMST}(T)$. This requires $O(N_3\rho^2 280^k)$ time to
 418 determine the optimal forest, where N_3 represents the number of non-empty clusters up to
 419 column 3. We store the backbone forest restricted to columns c_2 and c_3 along with the weight
 420 of the minimum spanning forest, computed using the dynamic programming approach from
 421 Lemma 4. Notice that no BBST connect two clusters which are two columns (rows) apart
 422 from each other. Thus, we only need to look at two consecutive columns and enumerate
 423 possible forests as backbones.

424 Next, we apply a sliding window technique, moving from left to right with a window of
 425 size $k \times 2$. At step $i - 1$, we assume that the minimum spanning forest whose backbone
 426 belongs to $\mathcal{F}(i - 2, i - 1)$ has already been computed. To expand to column i , we must
 427 consider all possible forests in $\mathcal{F}(i - 1, i)$ that overlap with $\mathcal{F}(i - 2, i - 1)$. This requires
 428 evaluating forests in $\mathcal{F}(i - 2, i - 1, i)$ to identify the valid overlaps. The selection of actual
 429 points from each cluster is again determined using the dynamic programming approach from
 430 Lemma 4, which runs in $O(N_i\rho^2)$ time for each given forest, where N_i represents the number
 431 of non-empty cells across the three consecutive columns c_{i-2}, c_{i-1}, c_i . Since there are at most
 432 280^k such forests, the runtime at step i is $O(N_i\rho^2 280^k)$. Thus, upon reaching the last column,
 433 the overall running time is $O(N\rho^2 280^k)$. Notice that at the end of the last row, we need to
 434 make sure we end up with a tree, eliminating the forests that are not tree after all. ◀

435 ▶ **Corollary 15.** *There is $(1 + \epsilon)$ -approximation algorithm for GGMST of \mathcal{I} in \mathbb{R}^2 that runs
 436 in $O(N\rho^2 280^{\frac{1}{\epsilon}})$ time.*

437 **Proof.** We apply the approximation algorithm from [5] to obtain an $(1 + \epsilon)$ -approximation.
 438 Their approach includes a subroutine that computes OptMST within a grid of size $k \times l$,
 439 running in time $O(l\rho^{6k} 2^{34k^2} k^2)$. However, as shown in Theorem 14, this step can be performed
 440 in $O(\rho^2 N 280^k)$ time. By carefully analyzing their PTAS, we conclude that a GMST T can
 441 be computed such that its weight is at most $(1 + \epsilon) \cdot \text{OptMST}$ while maintaining an overall
 442 runtime of $O(N\rho^2 280^{\frac{1}{\epsilon}})$. ◀

443 ▶ **Definition 16 (γ -bounded).** *Let $(\mathcal{G}, \mathcal{N}, P)$ be an instance of the GGMST. We say $(\mathcal{G}, \mathcal{N}, P)$
 444 is γ -bounded if for every non-empty cluster (i, j) , (in row i and column j), the number of
 445 non-empty clusters in row i is at most γ or the number of non-empty clusters in column j is
 446 at most γ .*

447 By a simple adaptation of the argument of the Lemma 13 we have the following lemma.

448 ▶ **Lemma 17.** *The number of backbone forests for instance $(\mathcal{G}, \mathcal{N}, P)$ which lies in a $l \times 3$
 449 grids and for each column 1, 2, 3 have at most γ non-empty cluster is at most 280^γ .*

23:12 Approximate and Exact Geometric Generalized Minimum Spanning Trees

	c_1	$c_1 + 1$	$c_1 + 2$		c_2	$c_2 + 1$	$c_2 + 2$	
	Empty		..		Empty	..		Empty
r_1
$r_1 + 1$
$r_1 + 2$
	Empty		Empty	Empty
r_2				
$r_2 + 1$
$r_2 + 2$

■ **Figure 2** Example of the γ -bounded instance. Note that there could be non-empty cluster in row c_1 as well as in row r_1 , etc. The empty text means the clusters without any point in them.

450 By a simple adaptation of the argument of the Theorem 14 we have the following lemma.

451 ► **Lemma 18.** *Let $(\mathcal{G}, \mathcal{N}, P)$ be an instance of the GGMST problem so that the number of
452 non-empty clusters in every column is at most γ . Then optimal GGMST can be calculated
453 with running time $O(N\rho^2 280^\gamma)$.*

454 We obtain the following theorem whose proof is delivered in Appendix.

455 ► **Theorem 19.** *Let $\mathcal{I} = (\mathcal{G}, \mathcal{N}, P)$ be an instance of GGMST in \mathbb{R}^2 that is γ -bounded with
456 $|\mathcal{N}| = N$ and $|P| = n$. The GGMST of \mathcal{I} can be computed in $O(N\rho^2 560^\gamma)$ time.*

457 6 Concluding remarks

458 In this paper we proposed some exact and approximation algorithms improving the runtime
459 on the latest results on the GGMST problem. Our algorithms work in any dimension d and
460 make use of some known techniques, the one by Arora appeared in [2]. Our main contribution
461 is indeed the $(1 + \epsilon)$ -approximation for GGMST in \mathbb{R}^d . Arora's original PTAS for TSP
462 relies on portal pairings—ensuring entry and exit in matching pairs—to construct a tour,
463 whereas our GMST variant instead tracks connected subsets (partitions) of portals to capture
464 how GMST edges merge boundary portals into forests, achieving an approximation error
465 of $\epsilon \cdot W(T_{opt})$ via a geometric series and Bell number analysis. Unlike TSP, where every
466 point must appear in the tour, our formulation requires exactly one point per cluster; these
467 activations are encoded at the leaf cells of our dynamic program to prevent redundancy at
468 higher levels. Furthermore, since edges can only connect adjacent unit squares, the complexity
469 associated with long-range edges is reduced. While Arora's method typically places $O(\frac{\log n}{\epsilon})$
470 portals per cell side—resulting in runtime factors such as $n^{\mathcal{O}(1/\epsilon)}$ —our approach utilizes
471 only $O(\frac{1}{\epsilon})$ portals, distributing overhead hierarchically across the quadtree levels, thereby
472 avoiding the extra $\log n$ factor in portal spacing. As with Arora's technique, the overall
473 runtime depends super-polynomially on $\frac{1}{\epsilon}$, confirming that our method is a PTAS: it is
474 polynomial in n for any fixed ϵ , but not fully polynomial in $\frac{1}{\epsilon}$.

475 Our other algorithms are also improving on some known results by [5] in \mathbb{R}^2 and provide
476 fixed parameter tractability for the case where the GMST instances behave naturally such
477 as having bounded treewidth and/or γ -bounded as we introduced in this paper.

478 ————— **References** —————

- 479 1 Steen Arnborg, Derek G. Corneil, and Andrew Proskurowski. Complexity of finding embeddings
480 in a k-tree. *SIAM Journal on Algebraic and Discrete Methods*, 8(2):277–284, 1987.
- 481 2 Sanjeev Arora. Polynomial-time approximation schemes for euclidean tsp and other geometric
482 problems. *Journal of the ACM*, 45(5):753–782, 1998.
- 483 3 B. Bhattacharya, A. Ćustić, A. Rafiey, and V. Sokol. Approximation algorithms for
484 generalized mst and tsp in grid clusters. In Z. Lu, D. Kim, W. Wu, W. Li, and D. Z. Du, editors,
485 *Combinatorial Optimization and Applications*, volume 9486 of *Lecture Notes in Computer
486 Science*, pages 117–129. Springer, Cham, 2015. doi:10.1007/978-3-319-26626-8_9.
- 487 4 Binay K. Bhattacharya, Ante Custin, Akbar Rafiey, Arash Rafiey, and Vladislav Sokol.
488 Approximation algorithms for generalized MST and TSP in grid clusters. In *Combinatorial
489 Optimization and Applications - 9th International Conference, COCOA 2015, Houston, TX,
490 USA*, volume 9486 of *Lecture Notes in Computer Science*, pages 110–125. Springer, 2015.
491 doi:10.1007/978-3-319-26626-8_9.
- 492 5 C. Feremans, A. Grigoriev, and R. Sitters. The geometric generalized minimum spanning tree
493 problem with grid clustering. *4OR*, 4:319–329, 2006.
- 494 6 M. R. Garey and David S. Johnson. *Computers and Intractability: A Guide to the Theory of
495 NP-Completeness*. W. H. Freeman, 1979.
- 496 7 Bruce L. Golden, S. Raghavan, and Daliborka Stanojevic. Heuristic search for the generalized
497 minimum spanning tree problem. *INFORMS Journal on Computing*, 17(3):290–304, 2005.
- 498 8 Sariel Har-peled. *Geometric Approximation Algorithms*. American Mathematical Society, USA,
499 2011.
- 500 9 He Jiang and Yudong Chen. An efficient algorithm for generalized minimum spanning tree
501 problem. In *Proceedings of the Genetic and Evolutionary Computation Conference (GECCO)*,
502 pages 217–224, 2010.
- 503 10 H. A. Kachooei, M. Davoodi, and D. Tayebi. On the generalized minimum spanning tree in
504 the euclidean plane. In *Proceedings of the 1st Iranian Conference on Computational Geometry
505 (ICCG)*, pages 19–23, 2018.
- 506 11 Ton Kloks. *Treewidth: computations and approximations*, volume 842 of *Lecture Notes in
507 Computer Science*. Springer-Verlag, Berlin, 1994.
- 508 12 Tuukka Korhonen. A single-exponential time 2-approximation algorithm for treewidth. In
509 *Proceedings of the 62nd IEEE Annual Symposium on Foundations of Computer Science (FOCS)*,
510 2021.
- 511 13 Young-Soo Myung, Chang ho Lee, and Dong wan Tcha. On the generalized minimum spanning
512 tree problem. *Networks*, 26(4):231–241, 1995.
- 513 14 Petrica C. Pop, Georg Still, and Walter Kern. An approximation algorithm for the generalized
514 minimum spanning tree problem with bounded cluster size. In *Proceedings of the Algorithms
515 and Complexity in Discrete Mathematics (ACiD)*, pages 115–121, 2005.
- 516 15 Petrica Claudiu Pop. The generalized minimum spanning tree problem. 2002.

517 **7 Appendix: Omitted proofs and pseudocodes**

518 ▶ **Lemma 6.** Let T_{opt} be an GGMST of cost $W(T_{opt})$ in a bounding box of side-length M .
 519 Suppose at quadtree level i , each cell has side length $\frac{M}{2^i}$. Let X_i be the number of GGMST
 520 edges that cross cell boundaries at level i . If every crossing edge is at least $\alpha \frac{M}{2^i}$ in length,
 521 then:

522
$$X_i \leq \frac{\beta W(T_{opt})}{\alpha (M/2^i)}$$

523 for some constant $\alpha > 0$ and $\beta > 0$.

524 **Proof.** Any GGMST edge e crossing a cell boundary at level i has length at least $\alpha \frac{M}{2^i}$,
 525 because it spans distinct cells of side $\frac{M}{2^i}$, for some $\alpha > 0$. Suppose there are X_i such crossing
 526 edges. Then these X_i edges alone have total length

527
$$\sum_{e \text{ crosses at level } i} \|e\| \geq X_i \left(\alpha \frac{M}{2^i} \right).$$

528 Applying a standard geometric pigeonhole argument, the GGMST cannot afford a set of
 529 edges whose combined length exceeds $\beta \cdot W(T_{opt})$, for some constant $\beta > 0$. Concretely, if
 530 the total length of all crossing edges at level i were allowed to surpass $\beta \cdot W(T_{opt})$, one could
 531 reconstruct a substructure heavier than $W(T_{opt})$, contradicting T_{opt} 's minimality.

532 Consider the MST edges crossing cell boundaries at level i , each having length at least
 533 $\alpha \frac{M}{2^i}$. The *pigeons* here are precisely those boundary-crossing edges themselves, while the
 534 *holes* represent “slots” in the MST’s total cost $W(T_{opt})$. In other words, if each crossing
 535 edge (pigeon) demands at least an $\alpha \frac{M}{2^i}$ share of the MST length budget (a hole), and the
 536 MST cannot allocate more than $W(T_{opt})$ total, then having “too many” such edges (pigeons)
 537 would exceed the MST’s entire weight (holes). Hence:

538
$$\sum_{e \text{ crosses at level } i} \|e\| \leq \beta W(T_{opt}).$$

539 Combining both inequalities yields:

540
$$X_i \left(\alpha \frac{M}{2^i} \right) \leq \beta W(T_{opt}),$$

541 which rearranges to

542
$$X_i \leq \frac{\beta W(T_{opt})}{\alpha \left(\frac{M}{2^i} \right)} = \frac{2^i \beta W(T_{opt})}{\alpha M}.$$

543 If X_i were any larger, the total cost of those crossing edges alone would exceed $W(T_{opt})$,
 544 which is a contradiction. ◀

545 ▶ **Theorem 19.** Let $\mathcal{I} = (\mathcal{G}, \mathcal{N}, P)$ be an instance of GGMST in \mathbb{R}^2 that is γ -bounded with
 546 $|\mathcal{N}| = N$ and $|P| = n$. The GGMST of \mathcal{I} can be computed in $O(N\rho^2 560^\gamma)$ time.

547 **Proof.** Consider that \mathcal{G} is embedded within a grid of dimensions $\alpha \times \beta$. We partition this
 548 grid as follows : Let $c_1, c_1 + 1, \dots, c_1 + \ell_1$ denote a sequence of consecutive columns where
 549 each column contains at most γ non-empty clusters. Similarly, let $c_2, c_2 + 1, \dots, c_2 + \ell_2$ be the
 550 next sequence of consecutive columns satisfying where each of them has at most γ non-empty
 551 clusters. We continue this partitioning until reaching the last column β (refer to Figure 2
 552 for visualization). Likewise, define $r_1, r_1 + 1, \dots, r_1 + s_1$ as the first set of consecutive rows

553 where each row contains at most γ non-empty clusters. Define $r_2, r_2 + 1, \dots, r_2 + s_2$ as the
 554 next set of consecutive rows where each also contains at most γ non-empty clusters. We
 555 continue this partitioning until reaching the last row α .

556 Using Lemma 18, we can compute the minimum spanning forest for sub-grid restricted
 557 to columns from c_1 to $c_1 + \ell_1$ and all the α rows. Next, we compute the optimal minimum
 558 forest within the rows $r_1, r_1 + 1, \dots, r_1 + s_1$ up to column $c_1 + 1$, denoted as $\text{OptMST}(F_{r_1})$.
 559 Observe that the number of possible forests in these rows and the corresponding column c_1 is
 560 bounded by 280^γ . Since the dynamic programming approach processes the rows sequentially,
 561 we ensure that the constructed forest is compatible with the existing forest on column c_1 .
 562 This requires $O(N'\rho^2 280^\gamma \times 2^\gamma)$ time, where N' represents the number of non-empty clusters
 563 in the rows from r_1 to $r_1 + s_1$. Observe that there are at most 2^γ possible forests on column
 564 c_1 . A similar procedure is applied to the subsequent set of rows, $r_2, r_2 + 1, \dots, r_2 + s_2$, and
 565 the next block of columns $c_2, c_2 + 1, \dots, c_2 + s_2$. To track all possible forests restricted to
 566 any row and up to column $c_1 + 1$, we spend $O(N_1\rho^2 280^\gamma \times 2^\gamma)$ time, where N_1 represents
 567 the number of non-empty clusters in this restricted section. Finally, as we process the
 568 columns $c_1 + 1, c_1 + 2, \dots, c_1 + \ell_1$, we maintain a sliding window of size $\alpha \times 2$. Note that for
 569 each of these forests there is another super-forest \mathcal{S} that is a potential part of the OptMST.
 570 We compute the $\text{OptMST}(\mathcal{S})$ using the dynamic programming approach in Lemma 4. The
 571 approach proceeds iteratively:

- 572 1. Process the grid portion between columns $c_1 + \ell_1$ and c_2 row by row until reaching c_2 .
- 573 2. Continue processing column by column.
- 574 3. Repeat 1 and 2 until we reach to the last row and last column in \mathcal{G} .

575 Thus, the overall runtime complexity of the algorithm is bounded by $O(N\rho^2 560^\gamma)$, as
 576 required. ◀

577 Algorithm Pseudocodes in Section 4

■ Algorithm 1 $\text{EnumerateSpanningTrees}(\mathcal{G})$

-
- 1: Compute an approximate treewidth k' and a tree decomposition $(T, \{X_t\})$ of \mathcal{G} using
 Korhonen's algorithm.
 - 2: Convert $(T, \{X_t\})$ into a nice tree decomposition.
 - 3: **for** each node t in post-order traversal of T **do**
 - 4: **if** t is a leaf node **then**
 - 5: $\text{DP}_t \leftarrow \{(\emptyset, \emptyset)\}$
 - 6: **else if** t is an introduce node introducing v **then**
 - 7: $\text{DP}_t \leftarrow \text{result of } \text{IntroduceNode}(t, v)$
 - 8: **else if** t is a forget node forgetting v **then**
 - 9: $\text{DP}_t \leftarrow \text{result of } \text{ForgetNode}(t, v)$
 - 10: **else if** t is a join node **then**
 - 11: $\text{DP}_t \leftarrow \text{result of } \text{JoinNode}(t)$
 - 12: At the root node r , output the spanning trees represented by DP_r .
-

Algorithm 2 IntroduceNode(t, v)

```

1: Initialize  $\text{DP}_t \leftarrow \emptyset$ 
2: for each state  $(\sigma', \delta')$  in  $\text{DP}_{t'}$  do
3:   for each subset  $\delta_v \subseteq E_v$  do
4:      $\delta \leftarrow \delta' \cup \delta_v$ 
5:      $\sigma \leftarrow$  result of UpdatePartition $(\sigma', v, \delta_v)$ 
6:     if Valid $(\sigma, \delta)$  then
7:       Add  $(\sigma, \delta)$  to  $\text{DP}_t$ 
return  $\text{DP}_t$ 
```

Algorithm 3 ForgetNode(t, v)

```

1: Initialize  $\text{DP}_t \leftarrow \emptyset$ 
2: for each state  $(\sigma', \delta')$  in  $\text{DP}_{t'}$  do
3:    $\sigma \leftarrow \sigma'$  with  $v$  removed
4:   Add  $(\sigma, \delta')$  to  $\text{DP}_t$ 
return  $\text{DP}_t$ 
```

Algorithm 4 JoinNode(t)

```

1: Initialize  $\text{DP}_t \leftarrow \emptyset$ 
2: for each state  $(\sigma_1, \delta_1)$  in  $\text{DP}_{t_1}$  do
3:   for each state  $(\sigma_2, \delta_2)$  in  $\text{DP}_{t_2}$  do
4:     if  $\sigma_1 = \sigma_2$  then
5:        $\delta \leftarrow \delta_1 \cup \delta_2$ 
6:       if Valid $(\sigma_1, \delta)$  then
7:         Add  $(\sigma_1, \delta)$  to  $\text{DP}_t$ 
return  $\text{DP}_t$ 
```

Algorithm 5 UpdatePartition(σ', v, δ_v)

```

1: Add  $v$  as a new part in  $\sigma'$ 
2: for each edge  $(v, u) \in \delta_v$  do
3:   Merge the parts containing  $v$  and  $u$  in  $\sigma'$ 
return  $\sigma'$ 
```

Algorithm 6 Valid(σ, δ)

```

1: Build a graph  $G_\delta = (X_t, \delta)$ 
2: if  $G_\delta$  contains a cycle then return False
3: if The connected components of  $G_\delta$  do not match  $\sigma$  then return False
return True
```
